

MBS-raw: Physical Optics for Ice Crystals

Manual

2026-03-22

Contents

1	Overview	3
2	Quick Start	3
3	CLI Reference	3
3.1	Particle Definition	3
3.1.1	-p TYPE HEIGHT DIAMETER [EXTRA]	3
3.1.2	-pf FILENAME	3
3.1.3	-rs SIZE	3
3.1.4	-forced_convex / -forced_nonconvex	3
3.2	Physical Parameters	4
3.2.1	-w WAVELENGTH_UM	4
3.2.2	-ri N_REAL N_IMAG	4
3.2.3	-abs	4
3.2.4	-n MAX_REFLECTIONS	4
3.3	Method	4
3.4	Orientations	4
3.4.1	-sobol N (recommended)	4
3.4.2	-adaptive EPS	4
3.4.3	-sym B G	4
3.4.4	Other orientation modes	5
3.5	Scattering Grid	5
3.5.1	-grid T1 T2 N_PHI N_THETA	5
3.5.2	-auto_tgrid (recommended)	5
3.5.3	-auto_phi	5
3.5.4	-tgrid FILENAME	5
3.6	Multi-Size (-sizefile)	5
3.7	Performance Options	6
3.8	Output	6
4	Output Files and Console Log	6
4.1	Output files	6
4.1.1	DIRNAME.dat — Mueller matrix	6
4.1.2	DIRNAME\out.txt — run summary	7
4.1.3	DIRNAME_jones.dat — Jones matrices (with -jones)	7
4.1.4	particle_for_check.dat — particle geometry	7
4.1.5	beam_log.dat — beam diagnostics	8
4.2	Console output (stdout)	8
4.2.1	Auto mode output	8
4.2.2	Sobol mode output	8

4.2.3	Scattering efficiency (always printed)	8
4.3	Console output (stderr)	8
4.3.1	Memory diagnostic	8
4.3.2	Timing breakdown (at program exit)	8
5	Build	9
5.1	Intel / AVX-512	9
5.2	AMD EPYC 7H12 (Zen 2, SP3)	9
5.3	AMD Zen 4 (Ryzen 7000, EPYC Genoa)	9
5.4	Generic AVX2	9
6	Performance	9
6.1	Single-thread benchmarks	9
6.2	Key optimizations	9
6.3	OpenMP scaling	10
7	CUDA / FFT Recommendations for File Particles	10
7.1	-gpu -fft	10
7.2	MBS_SHARED_ORIENT_CHUNK=N	10
7.3	MBS_SHARED_PIPELINE=1	11
7.4	MBS_SHARED_BETA_GROUP=N	11
8	Bugfixes	11
8.1	Forward-direction Fresnel sign (2026-03-22)	11
9	Known Limitations	11
10	File Listing	12
11	Bugfix: 15× Mismatch (old vs new binary)	12
11.1	Root causes	12
11.2	Validation	12
12	New Features (March 2026)	13
12.1	Adaptive theta grid (-auto_tgrid EPS)	13
12.2	Flag priorities	13
12.3	Multi-size (-multigrid Dmin Dmax N)	14
12.4	Per-beta save (-save_betas)	14
12.5	Parallel Phase 1 (tracing)	14
12.6	Additional bugfixes	14
13	Performance: MBS-fast vs MBS-raw (original)	14
13.1	Architecture comparison	15
13.2	Speedup breakdown	15
13.3	Measured speedup	15
13.4	Accuracy	15

1 Overview

MBS-raw computes light scattering by non-spherical particles using the Physical Optics (PO) method. It traces geometric optics rays through the particle (reflections, refractions), then applies Kirchhoff diffraction to each output beam to obtain the far-field Mueller matrix.

Optimized version: $\sim 70\times$ faster than original. AVX-512/AVX2 SIMD, OpenMP parallel, Sobol quasi-random orientations. Builds for Intel, AMD Zen2 (EPYC 7H12), and AMD Zen4 (Ryzen 7000 / EPYC Genoa).

2 Quick Start

```
# Simple run: hex column, Sobol orientations, auto theta grid
mbs_po --po --sobol 1024 --auto_tgrid \
  -p 1 10 10 -w 0.532 --ri 1.31 0 -n 12 \
  --grid 0 180 48 180 --close

# Fully adaptive (auto everything)
mbs_po --po --adaptive 0.01 --auto_tgrid \
  -p 1 10 10 -w 0.532 --ri 1.31 0 -n 12 \
  --grid 0 180 48 180 --close
```

3 CLI Reference

All flags are grouped by function. Required flags have no default; optional flags are marked with their defaults.

3.1 Particle Definition

3.1.1 -p TYPE HEIGHT DIAMETER [EXTRA]

Type	Shape	Parameters
1	Hexagonal column	H = height, D = diameter
2	Bullet	H, D (cap angle 62° , auto)
3	Bullet rosette	H, D, [cap height]
4	Droxtal	arg3 = sup parameter
10	Concave hexagonal	H, D, concavity depth
12	Hexagonal aggregate	H, D, number of elements

Symmetry: auto-detected. Hex column: $\beta_{\text{sym}} = \pi/2$, $\gamma_{\text{sym}} = \pi/3 \rightarrow 12\times$ reduction.

3.1.2 -pf FILENAME

Load particle geometry from file. -p argument provides the filename.

3.1.3 -rs SIZE

Resize particle loaded with -pf to maximal dimension SIZE (μm).

3.1.4 -forced_convex / -forced_nonconvex

Force convex or non-convex tracing algorithm.

3.2 Physical Parameters

3.2.1 -w WAVELENGTH_UM

Wavelength in micrometers. Examples: `-w 0.532` (green), `-w 1.064` (NIR).

3.2.2 -ri N_REAL N_IMAG

Complex refractive index $m = n_{\text{re}} + i n_{\text{im}}$.

Important: Without `-abs`, imaginary part is forced to zero.

3.2.3 -abs

Enable absorption. Requires `-w`. Uses `DiffractInclineAbs`.

3.2.4 -n MAX_REFLECTIONS

Value	Use case	Notes
5–6	Forward scattering, C_{sca}	$< 1\%$ error
12	General purpose	$< 5\%$ backscatter error
20	Backscattering	$< 3\%$ error
25	Reference calculations	

3.3 Method

- `-po` — Physical Optics (Kirchhoff diffraction). **Required for PO.**
- `-go` — Geometrical Optics (no diffraction).
- `-incoh` — Incoherent summation ($|J|^2$ per beam, then sum).
- `-karczewski` — Karczewski polarization matrix (experimental, same M_{11}).
- `-jones` — Output Jones matrices to file.
- `-all` — Compute all trajectory groups.

3.4 Orientations

3.4.1 -sobol N (recommended)

N Sobol quasi-random orientations (power of 2). Auto-uses particle symmetry.

Convergence: $O((\log N)^2/N)$, much faster than Monte Carlo $O(1/\sqrt{N})$.

x	$C_{\text{sca}} (< 2\%)$	$M_{11}(180^\circ) (< 5\%)$
10–50	128	512
50–200	256	1024
200–1000	512	4096
> 1000	1024	8192+

3.4.2 -adaptive EPS

Auto-determine orientations. Starts at 256, doubles until C_{sca} converges to relative accuracy EPS.

3.4.3 -sym B G

Override particle symmetry: $\beta \in [0, \pi/B]$, $\gamma \in [0, 2\pi/G]$.

Example: `-sym 6 2` $\rightarrow \beta \in [0, 30^\circ]$, $\gamma \in [0, 180^\circ]$ (12 $\times$ reduction).

3.4.4 Other orientation modes

- `-random N_BETA N_GAMMA` — Uniform grid.
- `-montecarlo N` — Random orientations.
- `-fixed BETA GAMMA` — Single orientation (degrees).
- `-orientfile FILE` — Read (β, γ) pairs (radians) from file.
- `-b MIN MAX, -g MIN MAX` — Override ranges (degrees).

3.5 Scattering Grid

3.5.1 `-grid T1 T2 N_PHI N_THETA`

Uniform grid. Use $N_\phi \geq 48$; $N_\phi = 6$ causes 50–200% error.

3.5.2 `-auto_tgrid` (recommended)

Auto non-uniform θ grid with 5 adaptive zones:

Zone	Range	Step	Purpose
Fine	0 to $10 \times \text{peak}$	$\text{peak}/20$	Forward peak
Transition	to 20°	geometric $\times 1.3$	Smooth bridging
Medium	to 120°	0.5° ($x < 100$) or 1°	Side scattering
Side/back	to 175°	1°	Backscattering
Near-backscatter	to 180°	0.25°	LDR




Figure 1: Left: total N_θ and N_ϕ vs size parameter x . Right: zone breakdown of θ -grid points.

3.5.3 `-auto_phi`

Auto-select N_ϕ based on x : $N_\phi = \max(48, 6\lceil\sqrt{x}/1.5\rceil)$.

Recommended: use both `-auto_tgrid` `-auto_phi` together.

3.5.4 `-tgrid FILENAME`

Custom θ grid from file (degrees, one per line).

3.6 Multi-Size (`-sizefile`)

```
--sizefile sizes.txt # one x = pi*D/lambda per line
```




Figure 2: Step size $\Delta\theta$ vs scattering angle for different x . Small x has coarser steps (broad peak); large x needs fine steps near $\theta=0$.

Traces beams once, caches geometry, recomputes diffraction for each size. Output: `M_x10.dat`, `M_x20.dat`, etc.

3.7 Performance Options

- `-beam_cutoff EPS` — Skip negligible beams.
- `-shadow_off` — Disable shadow beam (debug).
- `-r RATIO` — Small beam restriction ratio (default 100).

3.8 Output

- `-o DIRNAME` — Output folder (default `M`).
- `-close` — Exit after calculation (**required for scripts**).
- `-log SECONDS` — Progress output interval.
- `-tr FILE` — Trajectory groups from file.
- `-gr` — Per-group output.

4 Output Files and Console Log

4.1 Output files

All output is written to the directory specified by `-o` (default `M`).

4.1.1 `DIRNAME.dat` — Mueller matrix

Space-separated, 18 columns, one header line + one row per θ bin.

ScAngle	2pi*dcos	M11	M12	M13	M14	M21	M22	M23	M24
		M31	M32	M33	M34	M41	M42	M43	M44

- `ScAngle` — scattering angle θ in degrees
- `2pi*dcos` — solid angle weight: $2\pi \Delta \cos\theta$. Integration: $C_{\text{sca}} = \sum_j M_{11}(\theta_j) \times (2\pi \Delta \cos\theta_j)$




Figure 3: Left: N_ϕ from auto_phi (blue) vs fixed values. Right: Q_{sca} error vs N_ϕ for different sizes—larger x needs more ϕ bins.

- M11...M44 — 16 Mueller matrix elements (phi-averaged). For randomly-oriented particles:

$$M_{13} = M_{14} = M_{23} = M_{24} = M_{31} = M_{32} = M_{41} = M_{42} \approx 0.$$

Useful derived quantities:

$$Q_{sca} = C_{sca}/A_{proj}$$

$$g = \frac{\sum M_{11} \cos \theta \cdot 2\pi d \cos \theta}{\sum M_{11} \cdot 2\pi d \cos \theta} \quad (\text{asymmetry parameter})$$

$$-M_{12}/M_{11} = \text{degree of linear polarization}$$

4.1.2 DIRNAME\out.txt — run summary

Contains particle parameters, refractive index, method, start/end time, total number of orientations. Written after computation completes.

```
Particle: Refractive index: 1.310000
Area: 429.903838
Number of secondary reflections: 6
Wavelength (um): 0.532000
Method: Physical optics, Sobol quasi-random
Start of calc.: Sun Mar 22 11:22:44 2026
End of calc.: Sun Mar 22 11:22:49 2026
Total time : 4s
Total number of body orientations: 128
```

4.1.3 DIRNAME_jones.dat — Jones matrices (with -jones)

Written when -jones flag is used (fixed orientation only). 10 columns:

```
# theta(deg) phi(deg) J00_re J00_im J01_re J01_im J10_re J10_im J11_re J11_im
```

One row per (θ, ϕ) grid point. Complex Jones matrix $\mathbf{J}$ relates incident to scattered field:

$$\begin{pmatrix} E_{\parallel} \\ E_{\perp} \end{pmatrix}_{sca} = \mathbf{J} \begin{pmatrix} E_{\parallel} \\ E_{\perp} \end{pmatrix}_{inc}.$$

4.1.4 particle_for_check.dat — particle geometry

Dumped by the program for verification. Contains vertex coordinates and facet definitions of the particle as used in the calculation.

4.1.5 beam_log.dat — beam diagnostics

Written for the first orientation (debug). Columns:

```
# beam_num nActs area opticalPath |Jones| cx cy cz nVertices
```

nActs = number of internal interactions, cx,cy,cz = beam exit direction. The shadow beam has nActs = large number and cz = -1.

4.2 Console output (stdout)

The program writes progress and results to stdout.

4.2.1 Auto mode output

```
Auto theta grid: x=83.5, peak_width=2.16 deg, 464 points
Auto phi: x=83.5, N_phi=48
Adaptive mode: target relative change = 0.01
  N=64 (+64 new), M11(180)=7.09, dM11=100%, Csca=92.9, dCsca=100%
  N=128 (+128 new), M11(180)=6.69, dM11=5.6%, Csca=93.0, dCsca=0.12%
  N=256 (+256 new), M11(180)=6.64, dM11=0.8%, Csca=92.9, dCsca=0.08%
Converged at N=256 (both C_sca and M11)
```

Convergence messages:

- Converged (both C_sca and M11) — best: both criteria met
- Converged (C_sca ok, M11 within X%) — relaxed: M11 within $5 \times \varepsilon$
- Converged (C_sca stable, M11 still X%) — $C_{sca} < \varepsilon/10$ after 4 iterations
- WARNING: Max orientations reached — did not converge, result approximate

4.2.2 Sobol mode output

```
Sobol: 1024 orientations, beta_sym=30 deg, gamma_sym=180 deg
Phase 1 (tracing + preprocessing): 0.46 s
Phase 2 (diffraction, OpenMP): 13.25 s
Total: 13.71 s
```

4.2.3 Scattering efficiency (always printed)

```
===== SCATTERING EFFICIENCY =====
C_sca = 160.0398
A_proj (incoming energy) = 93.0208
Q_sca = C_sca / A_proj = 1.7205
=====
```

If $Q_{sca} > 2.5$: WARNING: $Q_{sca} > 2$. PO overestimates at this size.

4.3 Console output (stderr)

4.3.1 Memory diagnostic

```
Memory: 25219 MB available, chunk=4096 orientations (1 chunks)
Adaptive: max orientations = 32768 (25219 MB available)
```

4.3.2 Timing breakdown (at program exit)

```
===== HandlerPO TIMING BREAKDOWN =====
HandleBeams calls: 448
HandleBeams total: 69.255 s
  Beam loop total: 66.039 s
  AddToMueller: 2.606 s
  Unaccounted: 63.433 s
=====
```

“Unaccounted” = the optimized inline loop (batched sincos + edge loop + rotate_jones). Not instrumented to avoid timing overhead.

5 Build

5.1 Intel / AVX-512

```
bash build.sh # -> bin/mbs_po
```

-O3 -march=native -mavx512f -mavx512dq -fopenmp. Requires GCC ≥ 9 .

5.2 AMD EPYC 7H12 (Zen 2, SP3)

```
bash build_epyc.sh # -> bin/mbs_po_epyc
bash build_epyc_clang.sh # -> bin/mbs_po_epyc_clang
```

-march=znver2. No AVX-512; fast_sincos_8x falls back to 2×fast_sincos_4x.

NUMA: OMP_PROC_BIND=close OMP_PLACES=cores OMP_NUM_THREADS=64

5.3 AMD Zen 4 (Ryzen 7000, EPYC Genoa)

```
bash build_zen4.sh # -> bin/mbs_po_zen4
```

-march=znver4 -mavx512f -mavx512dq -mavx512vl.

5.4 Generic AVX2

Remove -mavx512* from build.sh. Code auto-detects via `#ifdef __AVX512F__`.

6 Performance

6.1 Single-thread benchmarks

Hex $D=H=10\mu\text{m}$, 128 Sobol, $48\phi \times 181\theta$, $n=6$:

Version	Phase 2	vs Original
Original (before optimizations)	~320 s	1×
Optimized + batched sincos	4.6 s	~70×
EPYC build (AVX2 only)	4.7 s	~68×

6.2 Key optimizations

1. **PrecomputeEdgeData**: slopes, intercepts computed once per beam (not per direction).
2. **Theta-coefficients**: $A(\theta) = \sin \theta \cdot a_s + \cos \theta \cdot a_c + a_0$ (2 FMA).
3. **Batched sincos**: all vertex phases for all θ pre-computed before the θ -loop. AVX-512 fast_sincos_8x or AVX2 2×fast_sincos_4x.
4. **Polynomial sincos**: Cody–Waite range reduction + Chebyshev polynomial (~ 14 decimal digits).
5. **Inline hot loop**: no heap allocation, stack-only complex arithmetic.
6. **OpenMP**: thread-local Mueller accumulators, `schedule(dynamic,1)`.
7. **Sobol + symmetry**: native C++ Sobol generator with particle symmetry reduction.

6.3 OpenMP scaling

Threads	Phase 2	Speedup
1	4.6 s	1.0×
4	1.7 s	2.7×
12 (SMT)	~0.9 s	~5×
64 (EPYC)	~0.2 s*	~23×

*Estimated; use ≥ 512 orientations for 64 cores.

7 CUDA / FFT Recommendations for File Particles

For large file particles and shared size sweeps, use the CUDA diffraction backend with FFT phi interpolation and the global orientation scheduler:

```
MBS_SHARED_PIPELINE=1 MBS_SHARED_ORIENT_CHUNK=64 \
./bin/mbs_po_float_fast --pf Afine30.dat --multikeq_list keq_list.txt \
  --ri 1.6 0.002 -n 14 --po --oldauto 2 \
  --beam_cutoff_j 0.001 --beam_cutoff_area 0.002 \
  --trace_cutoff_importance 0.0001 --trace_max_beams 20000 \
  --grid 0 180 600 181 -w 1.064 --gpu --fft --close
```

7.1 -gpu -fft

-gpu moves the Phase 2 diffraction integral to CUDA. The CPU still performs ray tracing and prepares beam data. **-fft** enables angular Fourier interpolation in φ : diffraction is evaluated on a reduced direct azimuth grid and reconstructed to the requested N_φ by cuFFT. This is not aperture pFFT/FMM and not a 3D spatial FFT.

Recommended scattering grid for ADDA comparisons used here:

```
--grid 0 180 600 181
```

This gives 600 azimuth points and 181 theta points, including 0 and 180 degrees.

7.2 MBS_SHARED_ORIENT_CHUNK=N

Shared **-oldauto -multikeq*** GPU runs use a global orientation scheduler. Orientations are indexed as

$$i = i_\beta N_\gamma + i_\gamma.$$

The scheduler traces contiguous chunks of N orientations and immediately sends each completed chunk to CUDA diffraction. This avoids waiting for a whole beta block when one non-convex orientation is slow.

Default for GPU shared multikeq:

```
MBS_SHARED_ORIENT_CHUNK=64
```

Memory scales approximately as

$$M_{\text{host}} \sim N_{\text{chunk}} \langle N_{\text{beams/orientation}} \rangle,$$

instead of

$$M_{\text{host}} \sim N_{\text{betaGroup}} N_\gamma \langle N_{\text{beams/orientation}} \rangle.$$

Practical choices:

- **MBS_SHARED_ORIENT_CHUNK=64**: recommended default for Afine30 / Greek file-particle runs.
- **MBS_SHARED_ORIENT_CHUNK=32**: lower latency; use when GPU utilization has long zero-utilization gaps.

- `MBS_SHARED_ORIENT_CHUNK=128`: fewer CUDA launches; test when GPU launch overhead dominates and host RAM is available.

On Afine30 with `-grid 0 180 600 181 -gpu -fft -oldauto 2`, chunks of 32–64 improved the observed scheduler rate from about 4.3 orientations/s to about 5.5 orientations/s on the RTX 4070 SUPER node and reduced host RAM from roughly 10 GB to below 2 GB.

7.3 MBS_SHARED_PIPELINE=1

This overlaps CPU tracing of chunk $k + 1$ with CUDA diffraction of chunk k . Use it with the global scheduler:

```
MBS_SHARED_PIPELINE=1 MBS_SHARED_ORIENT_CHUNK=64
```

If the GPU is still idle for long intervals, reduce the chunk to 32. If the GPU is busy but many tiny launches dominate, increase it to 128.

7.4 MBS_SHARED_BETA_GROUP=N

This is the older beta-block batching control. Large values reduce CUDA launch count but can be slower on non-convex particles because one hard orientation holds the whole beta group. With the global scheduler enabled, treat this as a diagnostic/legacy knob. Use:

```
MBS_SHARED_BETA_GROUP=1 # lowest latency
MBS_SHARED_BETA_GROUP=2 # conservative beta-block test
MBS_SHARED_BETA_GROUP=4 # larger packet, more RAM, longer tails
```

Recommended production mode is the global orientation chunk, not large beta groups.

8 Bugfixes

8.1 Forward-direction Fresnel sign (2026-03-22)

Bug: sign error in the forward-direction special case ($A \approx 0, B \approx 0$) of the Kirchhoff integral. The edge-sum converges to $-F_{\text{inv}} \cdot \text{area}$ but the shortcut used $+F_{\text{inv}} \cdot \text{area}$.

Symptom: $M_{11}(0^\circ)$ anomalously small in coherent mode for small particles. Example (hex $D=H=3\mu\text{m}$): $M_{11}(0)/M_{11}(0.1^\circ) = 0.019$ (should be ≈ 1.03).

Fix: negate the forward-direction formula in 6 locations:

- `HandlerPO.cpp`: lines 970, 1333, 1889
- `Handler.cpp`: lines 255, 473, 642

Incoherent mode unaffected ($|J|^2$ is sign-independent). Q_{sca} changed by $\sim 0.06\%$. See `BUGFIX_forward_direct` for full details.

9 Known Limitations

1. $Q_{\text{sca}} > 2$ at large x : PO does not enforce the optical theorem.
2. M_{33}, M_{34}, M_{44} : RotateJones vs Karczewski give different off-diagonal elements (same M_{11}).
3. **Small particles** ($x < 20$): PO not valid. Use ADDA or T-matrix.
4. **Backscattering convergence**: $M_{11}(180^\circ)$ needs 10–100 $\times$ more orientations.
5. **LTO**: `-flt0` causes $\sim 20\%$ regression. Do not use.

10 File Listing

File	Description
build.sh	Intel/AVX-512 build
build_epyc.sh	AMD EPYC 7H12 (Zen 2) build
build_epyc_clang.sh	EPYC + Clang/AOCC build
build_zen4.sh	AMD Zen 4 (Ryzen 7000/Genoa) build
Makefile	Legacy Makefile (Intel)
Makefile.epyc	Makefile for EPYC
tests/run_tests.sh	Regression test suite
MANUAL.md	This manual (Markdown source)
BUGFIX_forward_direction_sign.md	Bugfix documentation
docs/bugfix_15x_mismatch.pdf	15× mismatch investigation

11 Bugfix: 15× Mismatch (old vs new binary)

Comparison of `-fixed 45 30` between old and new binaries showed a $15.5\times$ difference in M_{11} .

11.1 Root causes

1. **Old `-fixed` never rotated the particle.** `ScatteringConvex::ScatterLight()` has `m_particle->Rotate` commented out, and old `TracerP0::TraceFixed` did not call it either. Result: `-fixed 45 30` always gave the default orientation. New code correctly adds `m_particle->Rotate(b, g, 0)` in `TraceFixed`.
2. **Uninitialized `d_param` in particle constructors.** `Point3f()` does not zero `coordinates[3]`. No constructor called `SetDParams()` after `Reset()`. Fixed: `SetDParams()` added after `Reset()` in all 8 particle constructors.
3. **Mixed `fast_exp_im/exp_im` in fallback paths.** `DiffractionIncline`, `DiffractionInclineFast`, `ComputeFnJones`, and `ApplyDiffraction*` had inconsistent precision. Fixed: all fallback paths use `exp_im` (glibc sincos).

11.2 Validation

Test	Max diff	Mode	Notes
<code>-fixed 0 0 -incoh</code>	exact	incoh	per-beam $ J ^2$ identical
<code>-random 64 43 -incoh</code>	exact	incoh	per-beam physics identical
<code>-random 32 22 -coh_orient</code>	10^{-7}	coh	machine precision
<code>-random 64 43 (default)</code>	0.007%	incoh avg	coh vs incoh averaging
HandleBeams inline vs fallback	identical		SIMD optimizations verified

Speedup: $55\times$ (23 min $\rightarrow$ 25 sec at 2860 orientations, 4 threads).

`d_param` note: `d_param` is the plane equation parameter ($ax + by + cz + d = 0$) stored in `Point3f::coordinates[3]`. It was uninitialized before `Rotate()` call. `SetDParams()` now added to all particle constructors. This changes coherent single-orientation M_{11} by $\sim 7\%$ (phase-sensitive), but incoherent $|J|^2$ is identical. After orientation averaging the effect is $< 0.01\%$.

Impact: Only `-fixed` mode affected. Orientation averaging modes (`-random`, `-sobol`, `-adaptive`, `-auto`) always call `Rotate()` via `TracerP0Total`, so `d_param` is always initialized correctly.

See `docs/bugfix_15x_mismatch.pdf` and `tests/reference_test/` for full reports and plots.

12 New Features (March 2026)

12.1 Adaptive theta grid (-auto_tgrid EPS)

Replaces the static zone-based theta grid with recursive bisection. For each interval $[\theta_i, \theta_{i+1}]$: computes M_{11} at the midpoint, compares with linear interpolation. If relative error $> \text{EPS}$, splits the interval.

- Minimum step: $\lambda/(8D)$ (quarter-Nyquist)
 - Maximum depth: 8 levels
 - Probe orientations: div16 physics estimate (same as adaptive start)
 - Result: 2–3 $\times$ fewer theta points than static grid
- Usage: `-auto_tgrid 0.05` (5% tolerance). Works with `-sobol`, `-random`, `-auto`.

12.2 Flag priorities

Multiple flags can set the same parameter. Resolution order (highest priority first):

Theta grid (N_θ and theta values)

1. `-tgrid FILE` — explicit non-uniform grid from file
2. `-grid T1 T2 Nphi Ntheta` — uniform grid from command line
3. `-auto_tgrid EPS` — adaptive bisection (probes M_{11})
4. `-auto` — implies `-auto_tgrid 0.05` (bisection)
5. `-oldauto DIV` — static zone grid from `ApplyAutoThetaGrid`
6. Default: placeholder (1 theta point)

Phi grid (N_φ)

1. `-nphi N` — explicit override (highest priority)
2. `-grid T1 T2 Nphi Ntheta` — `Nphi` is the 3rd argument
3. `-auto_phi` — formula $N_\varphi = \min(360, x/5 + 48)$
4. `-auto` — implies `-auto_phi`
5. `-oldauto DIV` — same formula as `-auto_phi`
6. Default: 1 (useless without override)

Orientation method

Mutually exclusive (first matched wins):

1. `-fixed BETA GAMMA` — single orientation
2. `-oldauto DIV` — physics-based $\beta \times \gamma$ grid
3. `-random Nb Ng` — regular grid
4. `-montecarlo N` — random uniform
5. `-orientfile FILE` — from file
6. `-auto EPS` / `-autofull EPS` — adaptive Sobol
7. `-adaptive EPS` — adaptive Sobol (manual grid)
8. `-sobol N` — fixed-count Sobol

Beam cutoff

1. `-beam_cutoff EPS` — explicit threshold
2. `-auto/-adaptive` — uses their EPS as cutoff
3. Default: 0 (no cutoff)

Coherent/incoherent

1. Default: coherent (Jones summed across beams per orientation)
2. `-incoh`: incoherent (Mueller per beam)
3. `-coh_orient`: coherent across ALL orientations (legacy)

Note: `-grid T1 T2 Nphi Ntheta` — the 3rd value is N_φ (not N_θ !). This is the original MBS convention. Use `-nphi` to avoid confusion.

12.3 Multi-size (`-multigrid Dmin Dmax N`)

Generates N sizes from $D_{\min}$ to $D_{\max}$ in logarithmic scale. For each size: rescales particle via `Resize()`, runs full `TraceFromSobol` (same optimized path as standalone). Results bit-exact with separate runs.

```
# 10 sizes from D=50 to D=500, reference particle D=500
bin/mbs_po --po --sobol 1024 --multigrid 50 500 10 \
  -p 1 500 62.5 -w 0.532 --ri 1.31 0 -n 8 \
  --grid 0 180 360 180 --close -o results
```

Output: `results_x{N}.dat` for each size parameter $x = \pi D/\lambda$.

12.4 Per-beta save (`-save_betas`)

With `-random`: writes per-beta Mueller contribution and cumulative Mueller to `_betas/` subfolder. Useful for convergence analysis and crash recovery.

12.5 Parallel Phase 1 (tracing)

Beam tracing (Phase 1) is now parallelized with OpenMP. Each thread gets its own copy of `Particle` + `Scattering` via `CloneFor()`. Phase 1 speedup: $5.4\times$ on 12 threads.

12.6 Additional bugfixes

- `-montecarlo` missing handler config (`isCoh`, `m_wave`, `shadowOff`)
- `DiffractionControlPoints`: wrong M11 formula (missing `d01`, `d10` terms)
- `-karczewski` ignored in fast diffraction path
- `rdtsc` inline asm incorrect on x86-64
- malloc leak in `Particle::SetFromFile`
- OOM in `TraceAdaptive`: sub-chunked to 4096 max
- `-grid` conflicts with auto theta/phi resolved
- `-oldauto` N rounding: ceil instead of ODD (smoother size scans)
- `-point` crash from uninitialized handler pointer
- MPI missing `}` in `TraceAdaptive` convergence block

13 Performance: MBS-fast vs MBS-raw (original)

All benchmarks: hex column $L = 199.5$, $D = 24.9375\ \mu\text{m}$, $\lambda = 0.532$, $m = 1.31$.

The original MBS-raw code is **single-threaded**, uses standard `sin/cos` (glibc), computes the diffraction integral with full `exp_im` per edge, and has no beam filtering or adaptive grids.

13.1 Architecture comparison

Component	MBS-raw (original)	MBS-fast
Phase 1 (tracing)	1 thread, sequential	OpenMP, <code>CloneFor()</code> per thread
Phase 2 (diffraction)	1 thread, <code>exp_im</code> per edge	OpenMP, batched <code>fast_sincos</code> (AVX)
Edge integral	<code>ChangeCoordinateSystem</code> per vertex	Precomputed edges + slopes
A,B coefficients	Recomputed per (beam, θ , φ)	Theta-coefficients: hoist per φ
Vertex phases	$2 \times \text{exp_im}$ per edge per θ	$1 \times \text{fast_sincos}$ per vertex, batched
Beam filtering	None (all beams computed)	Two-threshold cutoff (skip 45–80%)
θ grid	Uniform or hand-tuned zones	Adaptive bisection (<code>-auto_tgrid</code>)
φ grid	Manual <code>-grid</code>	Auto: $N_\varphi = x/5 + 48$
Orientations	Regular $\beta \times \gamma$ grid	Sobol quasi-random + adaptive
Memory	Full arrays, no chunking	Chunked streaming (4096 max)

13.2 Speedup breakdown

Each optimization contributes a multiplicative factor. Combined effect depends on which phase dominates (Phase 2 for most cases).

Optimization	Where	Factor	Notes
<i>Phase 2 (diffraction) — typically 90–95% of total time:</i>			
OpenMP (thread-local Mueller)	Phase 2	$N_{\text{cores}} \times$	12 cores $\rightarrow$ 10–11 $\times$
Batched sincos (AVX-512 8-wide)	Phase 2	$\sim 2 \times$	polynomial approx, SIMD
Precomputed edge data	Phase 2	$\sim 1.5 \times$	slopes, intercepts cached
Theta-coefficients	Phase 2	$\sim 1.3 \times$	A,B decomposed per φ
Beam cutoff	Phase 1+2	$\sim 2\text{--}4 \times$	skip weak beams
Adaptive theta (bisection)	Grid	$\sim 2\text{--}3 \times$	fewer θ points
<i>Phase 1 (tracing) — typically 5–10% of total time:</i>			
OpenMP (<code>CloneFor</code>)	Phase 1	$N_{\text{cores}} \times$	12 cores $\rightarrow$ 5.4 $\times$
<i>Adaptive mode:</i>			
DiffractControlPoints	Convergence	$\sim 100 \times$	4 angles, not full grid
Incremental Sobol	Convergence	$\sim 3.5 \times$	reuse previous orientations
Physics-based start (div16)	Convergence	$\sim 2 \times$	skip initial low- N iterations

Combined estimate for 12 cores:

Phase 2: $10 \times$ (OMP) $\times 2 \times$ (AVX) $\times 1.5 \times$ (edges) $\times 1.3 \times$ (theta-coeff) $\times 2 \times$ (beam cutoff) $\approx 78 \times$. With adaptive theta ($2 \times$): $\approx 156 \times$.

13.3 Measured speedup

Test	Params	Old (1 thr)	New	Speedup	Accuracy
Forward 0–25°	$n=4$, 64×43 , 4 thr	17 min	16 sec	63 $\times$	0.007%
Forward 0–25°	$n=4$, 64×43 , 12 thr	17 min	11.5 sec	89 $\times$	0.007%
Backward 160–180°	$n=11$, 64×43 , 4 thr	47 min	27 sec	105 $\times$	0.8%
Backward 160–180°	$n=11$, 64×43 , 12 thr	47 min	17 sec	168 $\times$	0.8%
<code>-coh_orient</code>	32×22 , 4 thr	30 sec	0.5 sec	60 $\times$	10^{-7}
Phase 1 only	$n=8$, 8192, 12 thr	5.2 sec	0.96 sec	5.4 $\times$	exact

On 128-core EPYC cluster with MPI: $\sim 1000 \times +$ expected.

13.4 Accuracy

- `-coh_orient`: machine precision (10^{-7}) match with original

- `-incoh`: exact per-beam $|J|^2$ — physics identical
- Default mode: $< 0.01\%$ at convergence (coh vs incoh averaging)
- Forward peak ($\theta=0$): 0.39% from `d_param` fix (correct value)
- All optimizations verified: inline path = fallback path (bit-exact)